

Lightweight open source on-spacecraft machine learning with mlpack and F-Prime

Ryan R. Curtin
NumFOCUS, Inc.
ryan@ratml.org

Scott M. Gilliland
Georgia Institute of Technology
scott.gilliland@gatech.edu

Sterling L. Peet
Georgia Institute of Technology
sterling.peet@gatech.edu

ABSTRACT

Onboard machine learning is an increasingly important topic in spaceflight systems, as machine learning and AI models are increasingly used for data processing, anomaly detection, attitude control systems, and other predictive applications. However, successfully deploying a machine learning model in a spaceflight computing environment is fraught with difficulty, as spaceflight computing environments are low-resource, but typical machine learning software solutions are implemented in languages like Python, which require an interpreter and other heavyweight dependencies. It is therefore virtually mandatory for any on-spacecraft machine learning to be written with lightweight toolkits that minimize dependencies and resource usage.

We demonstrate a solution to this problem: via the use of open source C++-based scientific computing software, including the mlpack machine learning library and Armadillo linear algebra library, we provide the first implementation of a working on-spacecraft machine learning anomaly detection system inside the F-Prime flight software framework that successfully detects non-trivial anomalous events by using a kernel density estimate on telemetry data. Our prototype is validated on physical hardware, where we physically created several anomalous situations that were all successfully detected. The code is available publicly on Github and can be used as a reference project for deploying any complex machine learning pipeline to a spacecraft via the F-Prime framework.

1 INTRODUCTION

Modern spaceflight missions produce massive amounts of data. NASA missions are now generating over 100 TB of data each day,¹ and this presents a significant challenge for data processing. Increasingly, a solution to this problem is direct on-board machine learning to process the data (as opposed to downlinking the data for ground-based processing).²⁻⁴ But the computational resources available on low-power satellites, such as CubeSats⁵ and other similar craft such as extraplanetary rovers are quite limited.

Often, the computational platforms of satellites can be somewhat esoteric. Although ARM-based platforms are likely the most common, with use in missions like SpaceCube⁶ and DODONA,⁷ legacy platforms are also in wide usage. For instance,

the ArgoMoon mission⁸ and Roman Space Telescope⁹ use SPARC processor. The common BAE RAD750¹⁰ is a PowerPC processor and is used on the Perseverance and Curiosity rovers.⁹ Even platforms like MIPS are in use on missions like New Horizons.¹¹

A further constraint on spaceflight computing systems is the flight software framework being used to control them. Two common choices, both supported by NASA, are F-Prime¹² and core Flight System (cFS).¹³ cFS is written in C, adhering to the C99 standard. Although this allows for precise programmer control in an embedded context, it makes interoperability with non-C99 toolkits difficult. For similar reasons, F-Prime is written in C++, targeting the C++11 standard or newer. This choice comes with the same interoperability difficulties.

This variety of computing platforms, the general

lack of computing resources present on them, as well as the language limitations imposed by cFS and F-Prime present a huge barrier to practitioners wishing to perform machine learning directly on a spacecraft. Most widely-used machine learning software packages, such as `scikit-learn`,¹⁴ TensorFlow,¹⁵ or PyTorch¹⁶ provide interfaces in Python. This is problematic for lightweight deployment: Python requires a language interpreter, which is inefficient and can be orders of magnitude larger than storage resources permit.¹⁷ Even projects aimed to productionalize machine learning models present problems: efforts like XLA,¹⁸ TensorFlow Lite, and ExecuTorch for PyTorch are cumbersome to build and link against, as they use complex build systems and have many dependencies. Vendor-specific alternatives like ARM’s CMSIS-NN¹⁹ suffer from limited vendor support, limited functionality, and cannot be deployed on other hardware. Worse, vendor-supplied toolkits are often arbitrarily EOL’ed, as in the case of several Qualcomm and Intel toolkits.

Yet there exists a widely-used native C++ scientific computing ecosystem, made up of tools like Armadillo,²⁰ Eigen,²¹ and xTensor²² for linear algebra, mlpack,²³ dlib-ml,²⁴ and ensmal²⁵ for machine learning, and numerous other libraries for related tasks. Many of these tools are intentionally lightweight, header-only, and dependency-free; thus they can be easily cross-compiled to any target architecture and deployed without burdensome dependency requirements. These libraries have been deployed widely, including for other on-board space-flight applications.^{26,27}

In this paper, we put both of these pieces together to solve the problem: we demonstrate that mlpack can be easily used inside of F-Prime to build a lightweight machine learning anomaly detector that runs directly on-board the spacecraft without the need for ground-based intervention. Following Peet et al.,²⁸ we deployed our solution to a Pololu Romi robot, which serves as a surrogate of a small satellite or extraplanetary rover. Our anomaly detector framework was able to successfully detect a number of complex anomalous situations that are not trivially detected with, e.g., thresholding.

Although our purposes here are to detect telemetry anomalies, our solution can be trivially adapted to other common on-board machine learning tasks, such as computer vision and image classification,^{29,30} health monitoring,³¹ and other tasks. Our code is publicly available for use and adaptation at <https://github.com/SterlingPeet/mlpack-FPrime-robot>, licensed under the TODO license.

2 SOFTWARE OVERVIEW

Our flight control software is implemented in C++ as an F-Prime deployment that internally uses mlpack’s kernel density estimation technique for its anomaly detection functionality.

2.1 F-Prime

As mentioned earlier, F-Prime is a modular flight software framework emphasizing the reusability of individual parts of software across different missions. In F-Prime, each part of the mission is decomposed into discrete *Components* with well-defined interfaces; these serve as the fundamental building blocks for the mission software.

A single mission consists of a number of components connected in a specific way, called a *deployment*. Each F-Prime component’s interface is defined as a number of input and output ports, which are defined and connected to other components. This interface and connections are specified in the Fpp domain-specific language.³² For example, the `Svc.SystemResources` component, which monitors the state of the computing environment (e.g. RAM usage, CPU utilization), typically outputs its telemetry via an Fpp connection to the mission telemetry service, which will report the telemetry to the ground control systems.

The F-Prime framework handles all of the communication between components, and provides numerous default components and deployments. This means that for a user, implementing the mission software stack typically means taking several F-Prime-supplied off-the-shelf components and connecting them with a few custom components implemented specifically for the mission.

Once the components are implemented and the connections in the deployment are defined, an F-Prime deployment can easily be cross-compiled to any target platform, including traditional platforms like VxWorks on the RAD750, and newer platforms like ARM systems running lightweight Linux distributions (for instance, a Raspberry Pi).

Compilation and configuration is handled via CMake, which is likely the most widely-used C++ build system generator. F-Prime automatically generates the CMake configuration based on the deployment topology. A user who has implemented their flight software with F-Prime can typically build and run it with a few simple commands via the use of the helper `fprime-util` and `fprime-gds` programs:

```
# automatically configure CMake
$ fprime-util generate
```

```
# build project with CMake
$ fprime-util build

# run the software and display
# ground control interface
$ fprime-gds
```

F-Prime deployments can work even on very resource-constrained hardware; for instance, it has been successfully used with program memory size as small as 128 kB and RAM as small as 64 kB.

Further resources on F-Prime can be found in the original F-Prime paper,¹² and more details related to our deployment topology can be found in the previous work of Peet et al.,²⁸ whose work we have used as a starting point.

2.2 mlpack

mlpack is a lightweight header-only machine learning library written in C++ with a focus on portability and efficiency.²³ It is built on top of the ensmallen mathematical optimization library,²⁵ the cereal serialization library,³³ and the well-known Armadillo linear algebra library²⁰ for its linear algebra primitives. It can also make use of the more recent Bandicoot library³⁴ to provide support for GPU-accelerated linear algebra operations. All of these libraries take advantage of C++’s template metaprogramming support to provide significant runtime accelerations and size reductions.

Importantly, mlpack and all of its dependencies are header-only, with the sole exception of Armadillo, which must be linked against a BLAS/LAPACK library. This significantly eases the compilation process, as users simply need to ensure that the headers are available. Although users must supply a compiled BLAS/LAPACK library specific to the target architecture, the OpenBLAS library³⁵ has easy support for cross-compilation to a wide variety of targets (including all spacecraft architectures described in the introduction), and mlpack has a CMake-based auto-downloader that can obtain all of the header-only dependencies of mlpack and automatically compile or cross-compile OpenBLAS without requiring user intervention.

From a high level, using mlpack for machine learning purposes is generally isomorphic to the workflows used in Python-based toolkits such as `scikit-learn`. Linear algebra is performed with Armadillo (or Bandicoot), where the API is expressly meant to mirror MATLAB to make transition to C++ easier for researchers and practitioners. For

example, this simple code snippet creates two random uniform matrices of size 100×50 and computes a third as $Z = (X + Y)/2$.

```
// For more details on the Armadillo API, see
// https://arma.sourceforge.net/docs.html.
arma::mat X(100, 50, arma::fill::randu);
arma::mat Y(100, 50, arma::fill::randu);
arma::mat Z = (X + Y) / 2;
```

Typical classification and regression algorithms in mlpack implement a simple and unified API. For example, supposing we wish to build a decision tree on a dataset of observations X and a vector of labels Y , we can use the following code.

```
// This matrix holds the predictors.
arma::mat X;
// This vector holds the labels
arma::Row<size_t> Y;

// Create the decision tree, optionally
// specifying parameters.
mlpack::DecisionTree d(/* ... */);

// Train the decision tree.
d.Train(X, Y);
```

The `d.Classify()` function could then be used to predict the class of a new data point using the trained decision tree. There are also a number of hyperparameters that control the tree-building procedure that can be specified in the constructor, as well as other methods that can be used to inspect or use the tree in other ways.

Other algorithms in mlpack follow a very similar construct-train-use paradigm, popularized originally by `scikit-learn`³⁶ and roughly reflected in virtually every modern machine learning and data science library. Available algorithms in mlpack range from common decision trees and logistic regression to more specialized tree-based algorithms³⁷ including tree-based density estimation techniques,^{38,39} which is what we have chosen to use here for our on-spacecraft anomaly detector.

3 HARDWARE OVERVIEW

Building on Peet et al.,²⁸ we choose to use the inexpensive Pololu Romi robot (shown in Figure 1) as a robotic explorer platform due to its standard hardware: an ATMega32U4-based mainboard controlled by a Raspberry Pi via an I2c connection; in our case, we used a Raspberry Pi 3B+ which contains a Cortex-A53 processor at 1.4 GHz and 1GB of RAM.



Figure 1: Pololu Romi robot.

From a control perspective, the Romi contains two wheels whose speeds are programmable, a speaker, and 3 LEDs (which we do not use for our work here). From a telemetry perspective, the following signals are available on the Romi and can be captured for F-Prime telemetry:

- Wheel turn counters (can be used for distance tracking)
- Acceleration sensors (x/y/z axis)
- Gyrometers (x/y/z)
- Ambient temperature
- Battery level

This is in addition to the default F-Prime telemetry running on the Raspberry Pi, which can track available memory, memory used, and CPU utilization percentages.

Although the Romi seems unlikely to survive for very long if deployed as-is in a spaceflight environment, we find it to be a compelling demonstration platform, as it reflects typical qualities of real spaceflight computing platforms: it has limited computational power and memory, as well as multiple interconnected devices, sensors, and motors that enable (semi-)autonomous operation, and its main processor architecture (ARM) is commonly used in COTS spaceflight applications.

4 OPERATING CONCEPT

The operating concept of our robotic explorer envisions three distinct operation settings:

- **Stationary.** In this mode, the explorer does not use its wheels to move, and it passively collects environmental data. It is expected in this

mode that neither the external environment (measured via temperature, acceleration) nor the internal environment (CPU usage) changes significantly.

- **Straight-line.** In this mode, the explorer is always moving forward or backward at either a low velocity or high velocity (but *not* an intermediate velocity or stopped). It is expected that the measured rate of travel matches the controlled wheel speeds, and that the controlled and measured velocities are neither zero nor an intermediate speed.
- **Spinning.** In this mode, the explorer's left wheel speed is equivalent to the right wheel speed, but in opposite directions, so the explorer is stationary but rotating. The rate of rotation can be arbitrary, but the explorer is expected to remain stationary and rotate at the rate given by the wheel speed.

During normal operation, an anomaly detection system that tracks previous telemetry is always running, and should the explorer enter an operating region where the collected telemetry does not match telemetry collected from known-good operation, an error will be issued. In our implementation, we simply used the speaker to indicate anomalies audibly. (In a real spaceflight application, it is likely that either a louder speaker or some other non-acoustic anomaly reporting means would be desired.)

Commands sent to the explorer can set it in any of these three modes, and reset the anomaly detection system such that its recent collected telemetry is interpreted as corresponding to the explorer working as desired.

5 ARCHITECTURE

In order to support the three modes of operation desired for our robotic explorer, we architect F-Prime components as shown in Figure 2. Each component already provided by F-Prime is printed in green, and components developed specifically for our mission requirements are in yellow. A description of each component's responsibilities is below:

- **cmdDisp:**
- **Prm DB:**
- **Ctrl:** overall control; issues commands to put the explorer in one of the three operating modes.
- **MtrCtrl:** motor control; sets the motor speeds and computes odometry




Figure 2: F-Prime component topology for our robotic explorer.

- **AnomDet**: density-estimation anomaly detector implemented with mlpack; receives telemetry from TlmSplit
- **HWDrv**: Romi hardware driver; issues hardware commands via I2c to I2CDrv
- **IMUSen**: Romi IMU sensor driver; issues hardware commands via I2c to I2CDrv
- **I2CDrv**: low-level I2C driver
- **TlmSplit**: telemetry splitter; routes all telemetry *both* to default TlmChan and also AnomDet
- **TlmChan**: telemetry collector
- **Event**: event bus

Note that because F-Prime restricts outputs from components to be routed to only one destination component, we implemented a telemetry splitter (TlmSplit) with one input port for input telemetry and two output ports. This allows us to duplicate

all telemetry to both the default F-Prime telemetry system and our custom anomaly detector.

6 ANOMALY DETECTION

Historically, anomaly detection was performed on spacecraft via simple thresholds applied to telemetry values.⁴⁰ However, this simple approach suffers from many false positives, needs tedious and error-prone tuning of the thresholds, and cannot easily detect situations where each individual telemetry value is within tolerance, but combinations of values indicate a malfunction. (Consider, as we will later, the trivial case of a wheel’s rate of motion not matching the power applied to it.)

Of course, more recent efforts have used far more complex techniques to detect anomalies, including LSTMs (neural networks),⁴¹ convolutional neural networks,⁴² and kernel methods,⁴³ to name a few. Of note is that the majority of these efforts are applying machine learning on the ground, instead of

directly on the spacecraft.

With F-Prime and mlpack, each of the techniques referenced above (and many more) can be deployed and run directly on the spacecraft, presuming that the spacecraft has the computational power necessary to run the technique. For ‘traditional’ machine learning techniques like a decision tree this is not generally a problem, but neural-network based approaches such as LSTMs can have much higher computational requirements, depending on the size of the network being used.

Because our efforts here are to demonstrate a working prototype, we opt to keep our machine learning simple and easy to understand. We will use a kernel density estimation (KDE) approach to model the probability density of observed telemetry states. Then, the anomaly detection system will compute the probability of the current telemetry state, and if that probability is exceedingly low, then the system concludes that we are in an anomalous state and emits an audible error tone via the buzzer.

This approach is best described visually, as in Figure 3. Due to the limitations of the medium, the figure can only show the density estimate in two dimensions; however, as we are receiving numerous channels of telemetry from various components, our density estimator as implemented uses 31 dimensions of telemetry. As this is a demonstration, we manually tune the density limit used to classify as an anomaly. In a real-world application that used this approach, the density should be tuned by collecting large amounts of non-anomalous data and building an accurate model of the distribution.

Implementationally speaking, mlpack’s `KDE` class is used to model the distribution. Internally, mlpack uses tree-based algorithms³⁸ to provide an efficient estimate of the density for a given telemetry point. To keep the tree small on a low-resource spacecraft, we use a configurable moving window of historical telemetry to model the probability density.

7 EXPERIMENTS

8 CONCLUSION




Figure 3: Kernel density estimation. The observed telemetry points (blue) are used to model the distribution (also blue) via a convolution with a Gaussian kernel. Then, new telemetry points are compared with the distribution (red). Point q_1 is outside the distribution and may be considered an anomaly.

References

- [1] Jane J. Lee and Ian J. O'Neill. NASA Turns to the Cloud for Help With Next-Generation Earth Missions, 2021.
- [2] David Henriquez. High Performance Spaceflight Computing (HPSC). Technical report, Pasadena, CA: Jet Propulsion Laboratory, National Aeronautics and Space Administration, 2016.
- [3] Tamer Mekky Ahmed Habib. Artificial intelligence for spacecraft guidance, navigation, and control: a state-of-the-art. *Aerospace Systems*, 5(4):503–521, 2022.
- [4] Tyler Cody and Paula J. Dempsey. Application of Machine Learning to Rotorcraft Health Monitoring. Technical report, Glenn Research Center, National Aeronautics and Space Administration, 2017.
- [5] Armen Poghosyan and Alessandro Golkar. Cubesat evolution: Analyzing CubeSat capabilities for conducting science missions. *Progress in Aerospace Sciences*, 88:59–83, 2017.
- [6] Cody Brewer, Nicholas Franconi, Robin Ripley, Alessandro Geist, Travis Wise, Sebastian Sabogal, Gary Crum, Sabrena Heyward, and Christopher Wilson. NASA SpaceCube intelligent multi-purpose system for enabling remote sensing, communication, and navigation in mission architectures. In *Small Satellite Conference 2020*, number SSC20-VI-07, 2020.
- [7] D-Orbit deploys USC’s La Jument system cubesat to prove-out space artificial intelligence (AI) technologies, 2024.
- [8] Marco Lombardo, Marco Zannoni, Igor Gai, Luis Gomez Casajus, Edoardo Gramigna, Riccardo Lasagni Manghi, Paolo Tortora, Valerio Di Tana, Biagio Cotugno, Simone Simonetti, et al. Design and analysis of the cis-lunar navigation for the ArgoMoon cubesat mission. *Aerospace*, 9(11):659, 2022.
- [9] Justin Goodwill, Gary Crum, James MacKinnon, Cody Brewer, Michael Monaghan, Travis Wise, and Christopher Wilson. NASA SpaceCube edge TPU smallsat card for autonomous operations and onboard science-data analysis. In *Proceedings of the Small Satellite Conference*, number SSC21-VII-08. AIAA, 2021.

- [10] Richard W Berger, Devin Bayles, Ronald Brown, Scott Doyle, Abbas Kazemzadeh, Ken Knowles, David Moser, John Rodgers, Brian Saari, Dan Stanley, et al. The RAD750: A Radiation Hardened PowerPC Processor for High Performance Spaceborne Applications. In *2001 IEEE Aerospace Conference Proceedings*, volume 5, pages 2263–2272. IEEE, 2001.
- [11] David Y. Kusnierkiewicz, Chris B. Hersman, Yanping Guo, Sanae Kubota, and Joyce McDevitt. A description of the Pluto-bound New Horizons spacecraft. *Acta Astronautica*, 57(2-8):135–144, 2005.
- [12] Robert Bocchino, Timothy Canham, Garth Watney, Leonard Reder, and Jeffrey Levison. F prime: an open-source framework for small-scale flight software systems. In *32nd Annual AIAA/USU Conference on Small Satellites*, 2018.
- [13] David McComas. Nasa/gsfcs flight software core flight system. In *Flight Software Workshop Flight Workshop*, number GSFC. CPR. 7525.2013, 2012.
- [14] Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, Jake Vanderplas, Alexandre Passos, David Cournapeau, Matthieu Brucher, Matthieu Perrot, and Edouard Duchesnay. Scikit-learn: Machine Learning in Python. *The Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [15] M. Abadi, P. Barham, J. Chen, Z. Chen, A. Davis, J. Dean, M. Devin, S. Ghemawat, G. Irving, M. Isard, M. Kudlur, J. Levenberg, R. Monga, S. Moore, D.G. Murray, B. Steiner, P. Tucker, V. Vasudevan, P. Warden, M. Wicke, Y. Yu, and X. Zheng. TensorFlow: A system for large-scale machine learning. In *12th USENIX Symposium on Operating Systems Design and Implementation (OSDI 16)*, pages 265–283, 2016.
- [16] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Köpf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala. PyTorch: An Imperative Style, High-Performance Deep Learning Library, 2019.
- [17] Ryan R. Curtin. Lightweight, low-overhead, high-performance: machine learning directly in C++. In *PyData NYC 2023*, 2023.
- [18] Amit Sabne. XLA: Compiling machine learning for peak performance. *Google Res*, 2020.
- [19] ARM Limited. CMSIS-NN: Efficient Neural Network Kernels for Arm Cortex-M CPUs. <https://github.com/ARM-software/CMSIS-NN>, 2024.
- [20] Conrad Sanderson and Ryan R. Curtin. Armadillo: a template-based C++ library for linear algebra. *Journal of Open Source Software*, 1(2):26, 2016.
- [21] Gaël Guennebaud, Benoit Jacob, et al. Eigen. <https://eigen.tuxfamily.org>, Accessed June 2024.
- [22] Johan Mabille, Sylvain Corlay, Wolf Vollprecht, et al. xTensor: multi-dimensional arrays with broadcasting and lazy computing. <https://xtensor.readthedocs.io/>, Accessed June 2024.
- [23] Ryan R. Curtin, Marcus Edel, Omar Shrit and Shubham Agrawal, Suryoday Basak, James J. Balamuta and Ryan Birmingham, Kartik Dutt, Dirk Eddelbuettel and Rishabh Garg, Shikhar Jaiswal, Aakash Kaushik and Sangyeon Kim, Anjishnu Mukherjee, Nanubala Gnana Sai and Nippun Sharma, Yashwant Singh Parihar, and Roshan Swain and Conrad Sanderson. mlpack 4: a fast, header-only C++ machine learning library. *Journal of Open Source Software*, 8(82):5026, 2023.
- [24] Davis E King. Dlib-ml: A Machine Learning Toolkit. *Journal of Machine Learning Research*, 10:1755–1758, 2009.
- [25] Ryan R. Curtin, Marcus Edel, Rahul Ganesh Prabhu, Suryoday Basak, Zhihao Lou, and Conrad Sanderson. The ensmallen library for flexible numerical optimization. *Journal of Machine Learning Research*, 22(166):1–6, 2021.

- [26] Timothy M. Hackett, Sven G. Bilén, Paulo Victor R. Ferreira, Alexander M. Wyglinski, and Richard C. Reinhart. Implementation of a space communications cognitive engine. In *2017 Cognitive Communications for Aerospace Applications Workshop (CCAA)*, pages 1–7. IEEE, 2017.
- [27] Timothy M. Hackett, Sven G. Bilén, Paulo Victor R. Ferreira, Alexander M. Wyglinski, Richard C. Reinhart, and Dale J. Mortensen. Implementation and on-orbit testing results of a space communications cognitive engine. *IEEE Transactions on Cognitive Communications and Networking*, 4(4):825–842, 2018.
- [28] Sterling L Peet, Scott M Gilliland, and Jeffrey S Young. The framework makes the mission-an analytical comparison of two popular nasa open source flight software framework offerings. In *Proceedings of the Small Satellite Conference 2024*, 2024.
- [29] Gianluca Giuffrida, Luca Fanucci, Gabriele Meoni, Matej Batič, Léonie Buckley, Aubrey Dunne, Chris Van Dijk, Marco Esposito, John Hefe, Nathan Vercruyssen, et al. The Φ -Sat-1 mission: The first on-board deep neural network demonstrator for satellite earth observation. *IEEE Transactions on Geoscience and Remote Sensing*, 60:1–14, 2021.
- [30] Vasileios Leon, Panagiotis Minaidis, George Lentaris, and Dimitrios Soudris. Accelerating ai and computer vision for satellite pose estimation on the intel myriad x embedded soc. *Microprocessors and Microsystems*, 103:104947, 2023.
- [31] Al Volponi and Donald L. Simon. Enhanced self tuning on-board real-time model (estorm) for aircraft engine performance health tracking. Technical Report NASA/CR-2008-215272, National Aeronautics and Space Administration, 2008.
- [32] Robert L Bocchino, Jeffrey W Levison, and Michael D Starch. Fpp: A modeling language for f prime. In *2022 IEEE Aerospace Conference (AERO)*, pages 1–15. IEEE, 2022.
- [33] W. Shane Grant and Randolph Voorhies. cereal—a C++11 library for serialization. <http://uscilab.github.io/cereal>, 2017.
- [34] Ryan R Curtin, Marcus Edel, and Conrad Sanderson. Fast gpu linear algebra via compile time expression fusion. *arXiv preprint arXiv:2604.22242*, 2026.
- [35] Z. Xianyi, W. Qian, and Z. Chothia. OpenBLAS. <http://xianyi.github.io/OpenBLAS>, January 2021.
- [36] Lars Buitinck, Gilles Louppe, Mathieu Blondel, Fabian Pedregosa, Andreas Mueller, Olivier Grisel, Vlad Niculae, Peter Prettenhofer, Alexandre Gramfort, Jaques Grobler, et al. Api design for machine learning software: experiences from the scikit-learn project. In *European Conference on Machine Learning and Principles and Practices of Knowledge Discovery in Databases*, 2013.
- [37] Ryan Curtin, William March, Parikshit Ram, David Anderson, Alexander Gray, and Charles Isbell. Tree-independent dual-tree algorithms. In *International Conference on Machine Learning*, pages 1435–1443. PMLR, 2013.
- [38] Alexander G Gray and Andrew W Moore. Nonparametric density estimation: Toward computational tractability. In *Proceedings of the 2003 SIAM International Conference on Data Mining*, pages 203–211. SIAM, 2003.
- [39] Parikshit Ram and Alexander G Gray. Density estimation trees. In *Proceedings of the 17th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 627–635, 2011.
- [40] Sylvain Fuertes, Gilles Picart, Jean-Yves Tournet, Lotfi Chaari, André Ferrari, and Cédric Richard. Improving spacecraft health monitoring with automatic anomaly detection techniques. In *14th international conference on space operations*, page 2430, 2016.
- [41] Kyle Hundman, Valentino Constantinou, Christopher Laporte, Ian Colwell, and Tom Soderstrom. Detecting spacecraft anomalies using lstms and nonparametric dynamic thresholding. In *Proceedings of the 24th ACM SIGKDD international conference on knowledge discovery & data mining*, pages 387–395, 2018.

- [42] Liang Liu, Ling Tian, Zhao Kang, and Tianqi Wan. Spacecraft anomaly detection with attention temporal convolution networks. *Neural Computing and Applications*, 35(13):9753–9761, 2023.
- [43] Ryohei Fujimaki, Takehisa Yairi, and Kazuo Machida. An approach to spacecraft anomaly detection problem using kernel feature space. In *Proceedings of the eleventh ACM SIGKDD international conference on Knowledge discovery in data mining*, pages 401–410, 2005.